

# Problem1

## Report

*Student NO: 20183784*  
*Student Name: 노현진*

[Execution environment]

- OS

DESKTOP-9RV3QK4  
Precision 3630 Tower

이 PC의 이름 바꾸기

i

장치 사양

복사

^

장치 이름

DESKTOP-9RV3QK4

프로세서

Intel(R) Core(TM) i9-9900K CPU @ 3.60GHz 3.60 GHz

설치된 RAM

32.0GB(31.8GB 사용 가능)

장치 ID

F89FE21E-99F5-4AEC-9F65-030082496F71

제품 ID

00330-52958-55864-AAOEM

시스템 종류

64비트 운영 체제, x64 기반 프로세서

펜 및 터치

이 디스플레이에 사용할 수 있는 펜 또는 터치식 입력이 없습니다.

관련 링크

도메인 또는 작업 영역

시스템 보호

고급 시스템 설정

Windows

Windows 사양

복사

^

에디션

Windows 11 Pro

버전

22H2

설치 날짜

2023-03-08

OS 빌드

22621.1702

경험

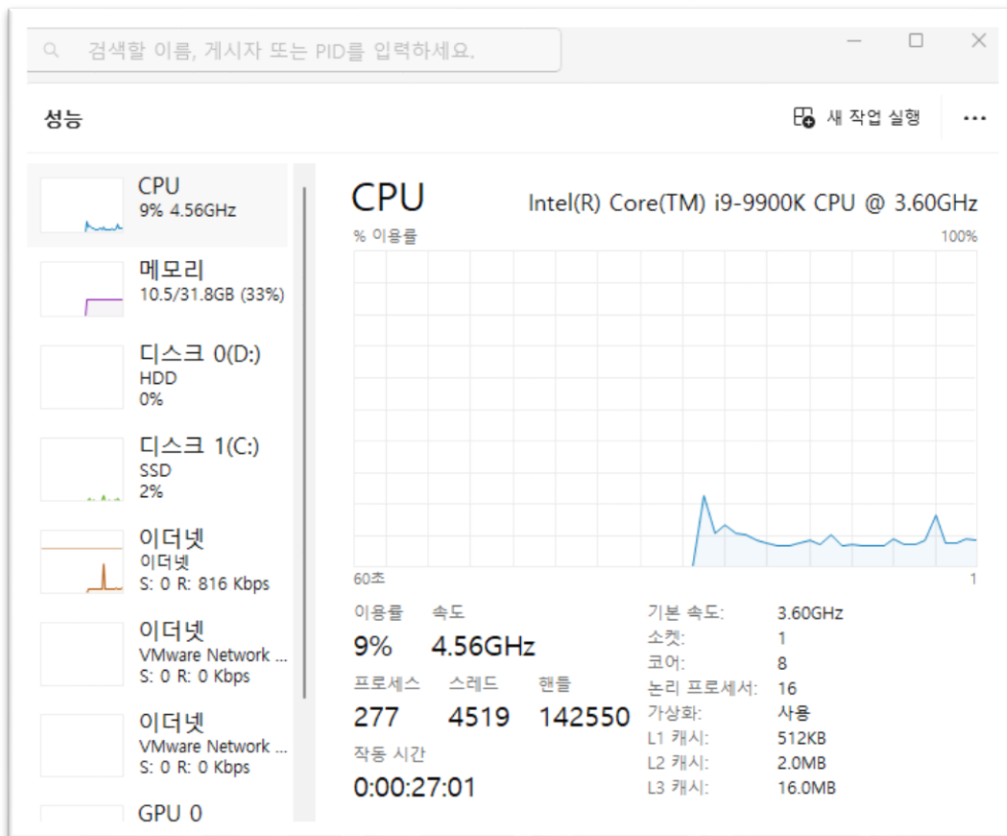
Windows Feature Experience Pack 1000.22641.1000.0

Microsoft 서비스 계약

Microsoft 소프트웨어 사용 조건

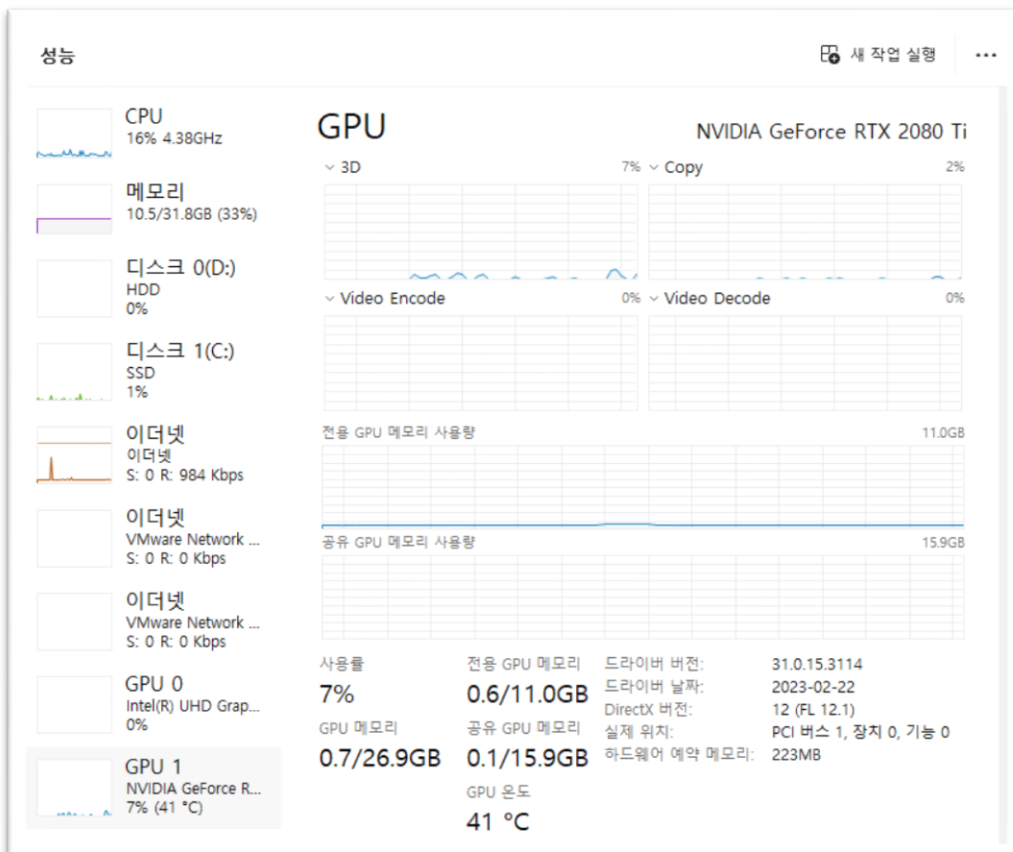
OS type: Windows 11

## CPU



**CPU type: Intel® Core™ i9-9900K CPU**

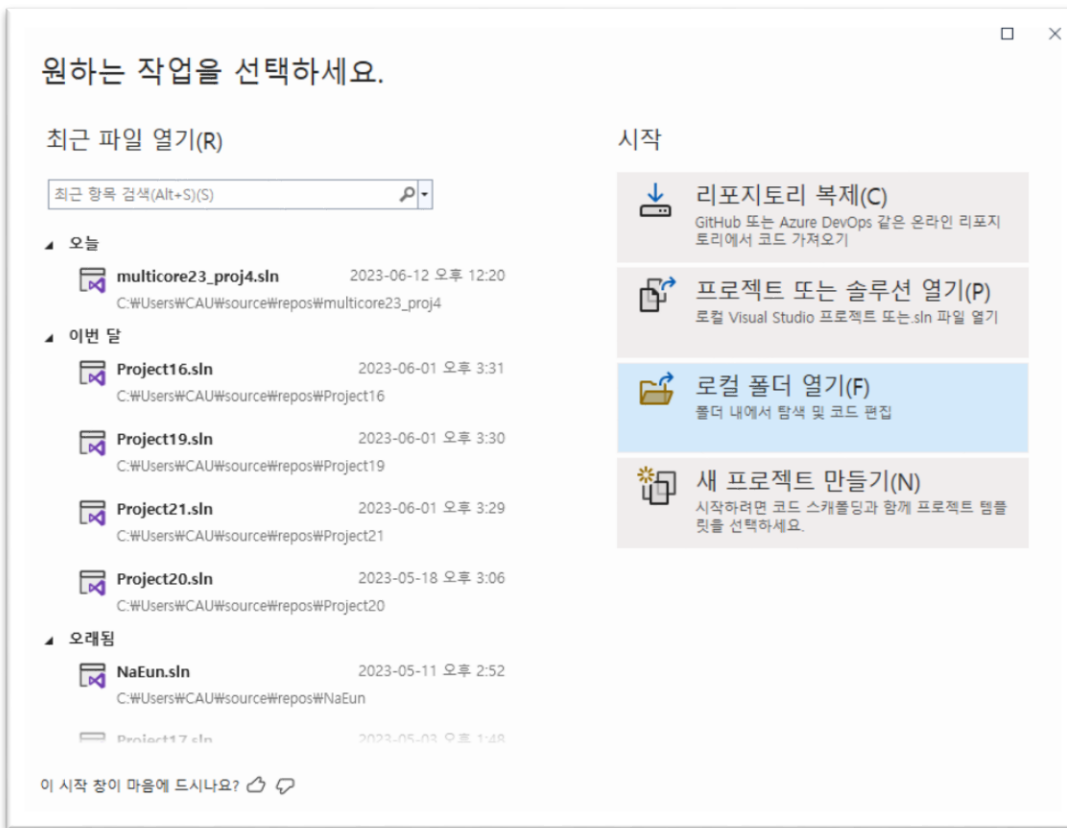
## GPU



**GPU type: NVIDIA GeForce RTX 2080 Ti**

## [How to compile and execute]

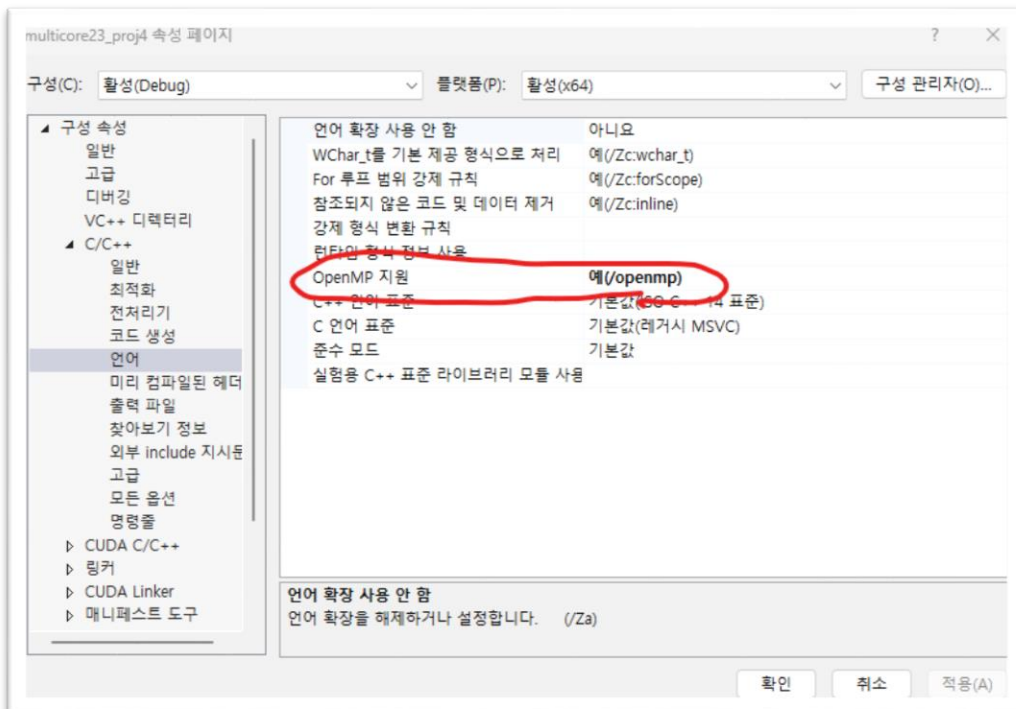
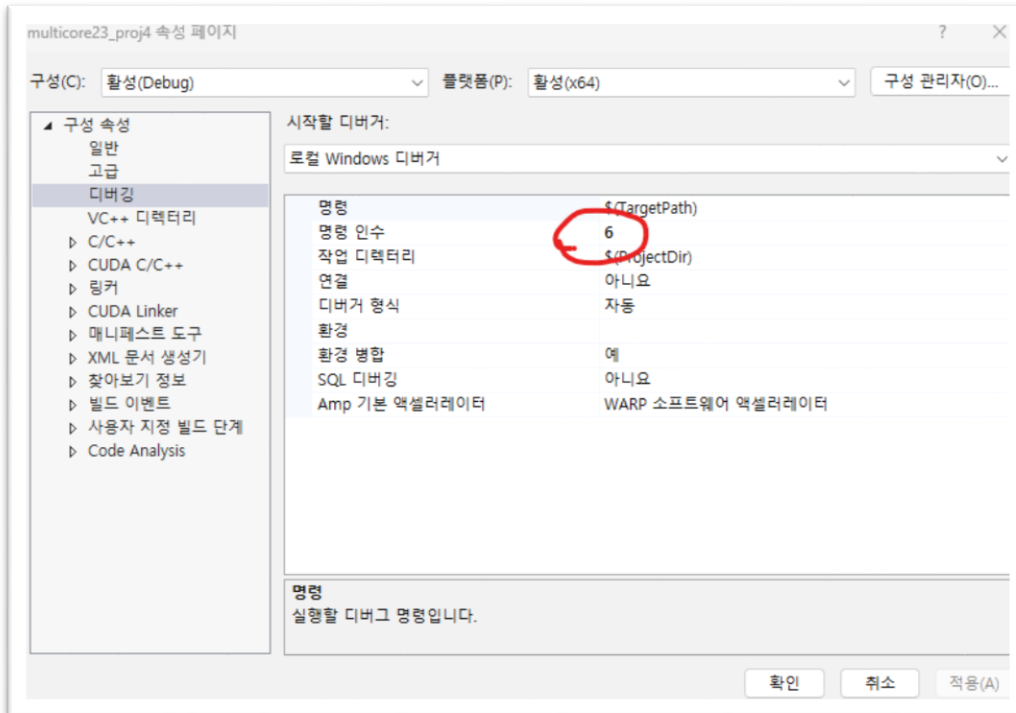
1. Firstly, install Visual Studio 2022.
2. Secondly, install CUDA.
3. After installations, open the Visual Studio 2022.



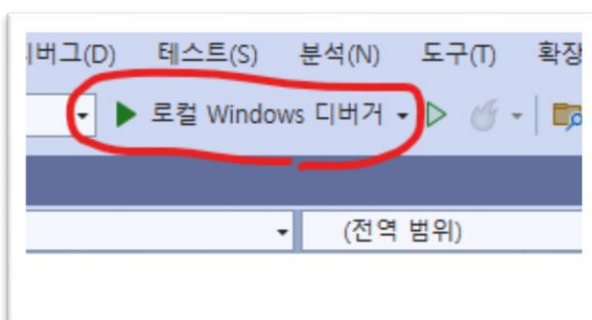
4. Create new project. The project should be CUDA Run time project.



5. Put the source code that you want to run into the project.
6. Before execution, set the parameter and make OpenMP available.



7. Finally, run the source code.



## [Source Code]

### - openmp\_ray.cpp

```
#include <omp.h>
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <time.h>
#include <math.h>

#define SPHERES 20
#define rnd( x ) (x * rand() / RAND_MAX)
#define INF 2e10f
#define DIM 2048

struct Sphere {
    float r, b, g;
    float radius;
    float x, y, z;
    float hit(float ox, float oy, float* n) {
        float dx = ox - x;
        float dy = oy - y;
        if (dx * dx + dy * dy < radius * radius) {
            float dz = sqrtf(radius * radius - dx * dx - dy * dy);
            *n = dz / sqrtf(radius * radius);
            return dz + z;
        }
        return -INF;
    }
};

void kernel(int x, int y, Sphere* s, unsigned char* ptr) {
    int offset = x + y * DIM;
    float ox = (x - DIM / 2);
    float oy = (y - DIM / 2);

    //printf("x:%d, y:%d, ox:%f, oy:%f\n", x, y, ox, oy);

    float r = 0, g = 0, b = 0;
    float maxz = -INF;

    for (int i = 0; i < SPHERES; i++) {
        float n;
        float t = s[i].hit(ox, oy, &n);
        if (t > maxz) {
            float fscale = n;
            r = s[i].r * fscale;
            g = s[i].g * fscale;
            b = s[i].b * fscale;
            maxz = t;
        }
    }

    ptr[offset * 4 + 0] = (int)(r * 255);
    ptr[offset * 4 + 1] = (int)(g * 255);
    ptr[offset * 4 + 2] = (int)(b * 255);
    ptr[offset * 4 + 3] = 255;
}
```

```

void ppm_write(unsigned char* bitmap, int xdim, int ydim, FILE* fp) {
    int i, x, y;
    fprintf(fp, "P3Wn");
    fprintf(fp, "%d %dWn", xdim, ydim);
    fprintf(fp, "255Wn");

    for (y = 0; y < ydim; y++) {
        for (x = 0; x < xdim; x++) {
            i = x + y * xdim;
            fprintf(fp, "%d %d %d ", bitmap[4 * i], bitmap[4 * i + 1], bitmap[4 * i + 2]);
        }
        fprintf(fp, "Wn");
    }
}

int main(int argc, char* argv[]) {
    int i;
    int num_threads;
    int x, y;
    unsigned char* bitmap;
    double start_time, end_time;
    Sphere* temp_s = (Sphere*)malloc(sizeof(Sphere) * SPHERES);
    FILE* fp = fopen("result.ppm", "w");
    srand(time(NULL));

    if (argc != 2) {
        printf("> a.out [number of threads]Wn");
        printf("for example, '> a.out 8' means executing OpenMP with 8 threadsWn");
        exit(0);
    }
    else {
        num_threads = atoi(argv[1]);
    }

    omp_set_num_threads(num_threads);
    start_time = omp_get_wtime();

    for (i = 0; i < SPHERES; i++) {
        temp_s[i].r = rnd(1.0f);
        temp_s[i].g = rnd(1.0f);
        temp_s[i].b = rnd(1.0f);
        temp_s[i].x = rnd(2000.0f) - 1000;
        temp_s[i].y = rnd(2000.0f) - 1000;
        temp_s[i].z = rnd(2000.0f) - 1000;
        temp_s[i].radius = rnd(200.0f) + 40;
    }

    bitmap = (unsigned char*)malloc(sizeof(unsigned char) * DIM * DIM * 4);

    #pragma omp parallel for default(shared) private(x, y)
    for (x = 0; x < DIM; x++)
        for (y = 0; y < DIM; y++) kernel(x, y, temp_s, bitmap);

    end_time = omp_get_wtime();

    ppm_write(bitmap, DIM, DIM, fp);

    fclose(fp);
    free(bitmap);
    free(temp_s);
}

```

```

printf("OpenMP (%d threads) ray tracing: %lf sec\n", num_threads, end_time - start_time);
printf("[result.ppm] was generated.\n");
return 0;
}

```

## - cuda\_ray.cu

```

#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <time.h>
#include <math.h>
#include <chrono>
#include <iostream>
#include "cuda_runtime.h"
#include "device_launch_parameters.h"

using namespace std;
using namespace std::chrono;

#define SPHERES 20
#define rnd( x ) (x * rand() / RAND_MAX)
#define INF 2e10f
#define DIM 2048

struct Sphere {
    float r, b, g;
    float radius;
    float x, y, z;
    float hit(float ox, float oy, float* n) {
        float dx = ox - x;
        float dy = oy - y;
        if (dx * dx + dy * dy < radius * radius) {
            float dz = sqrtf(radius * radius - dx * dx - dy * dy);
            *n = dz / sqrtf(radius * radius);
            return dz + z;
        }
        return -INF;
    }
};

__global__ void kernel(Sphere* s, unsigned char* ptr) {
    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;
    int offset = x + y * DIM;
    float ox = (x - DIM / 2);
    float oy = (y - DIM / 2);

    //printf("x:%d, y:%d, ox:%f, oy:%f\n", x, y, ox, oy);

    int i;
    float r = 0, g = 0, b = 0;
    float maxz = -INF;

```



```

for (i = 0; i < SPHERES; i++) {
    float n;
    //float t = s[i].hit(ox, oy, &n);
    float t;
    float dx = ox - s[i].x;
    float dy = oy - s[i].y;
    float radius = s[i].radius;
    if (dx * dx + dy * dy < radius * radius) {
        float dz = sqrtf(radius * radius - dx * dx - dy * dy);
        n = dz / sqrtf(radius * radius);
        t = dz + s[i].z;
    }
    else {
        t = -INF;
    }
    if (t > maxz) {
        float fscale = n;
        r = s[i].r * fscale;
        g = s[i].g * fscale;
        b = s[i].b * fscale;
        maxz = t;
    }
}

ptr[offset * 4 + 0] = (int)(r * 255);
ptr[offset * 4 + 1] = (int)(g * 255);
ptr[offset * 4 + 2] = (int)(b * 255);
ptr[offset * 4 + 3] = 255;
}

void ppm_write(unsigned char* bitmap, int xdim, int ydim, FILE* fp) {
    int i, x, y;
    fprintf(fp, "P3\n");
    fprintf(fp, "%d %d\n", xdim, ydim);
    fprintf(fp, "255\n");

    for (y = 0; y < ydim; y++) {
        for (x = 0; x < xdim; x++) {
            i = x + y * xdim;
            fprintf(fp, "%d %d %d ", bitmap[4 * i], bitmap[4 * i + 1], bitmap[4 * i + 2]);
        }
        fprintf(fp, "\n");
    }
}

int main() {
    int i;
    Sphere* temp_s;
    Sphere* d_temp_s;
    unsigned char* bitmap;
    unsigned char* d_bitmap;
    int size1 = sizeof(Sphere) * SPHERES;
    int size2 = sizeof(unsigned char) * DIM * DIM * 4;
    FILE* fp = fopen("result.ppm", "w");
    srand(time(NULL));

    // Allocate space for device copies of temp_s and bitmap
    cudaMalloc((void**)&d_temp_s, size1);
    cudaMalloc((void**)&d_bitmap, size2);
}

```

```

// Allocate space for host copies of temp_s and bitmap
temp_s = (Sphere*)malloc(size1);
bitmap = (unsigned char*)malloc(size2);

// Setup initial values
for (i = 0; i < SPHERES; i++) {
    temp_s[i].r = rnd(1.0f);
    temp_s[i].g = rnd(1.0f);
    temp_s[i].b = rnd(1.0f);
    temp_s[i].x = rnd(2000.0f) - 1000;
    temp_s[i].y = rnd(2000.0f) - 1000;
    temp_s[i].z = rnd(2000.0f) - 1000;
    temp_s[i].radius = rnd(200.0f) + 40;
}

auto start_time = high_resolution_clock::now();

// Copy values to device
cudaMemcpy(d_temp_s, temp_s, size1, cudaMemcpyHostToDevice);

// Setup the execution configuration
dim3 dimBlock(32, 32, 1);
dim3 dimGrid(64, 64, 1);

kernel<<<dimGrid, dimBlock>>>(d_temp_s, d_bitmap);

// Copy result back to host
cudaMemcpy(bitmap, d_bitmap, size2, cudaMemcpyDeviceToHost);

auto end_time = high_resolution_clock::now();
auto duration = duration_cast<microseconds>(end_time - start_time);

ppm_write(bitmap, DIM, DIM, fp);

fclose(fp);
free(bitmap);
free(temp_s);
cudaFree(d_bitmap);
cudaFree(d_temp_s);

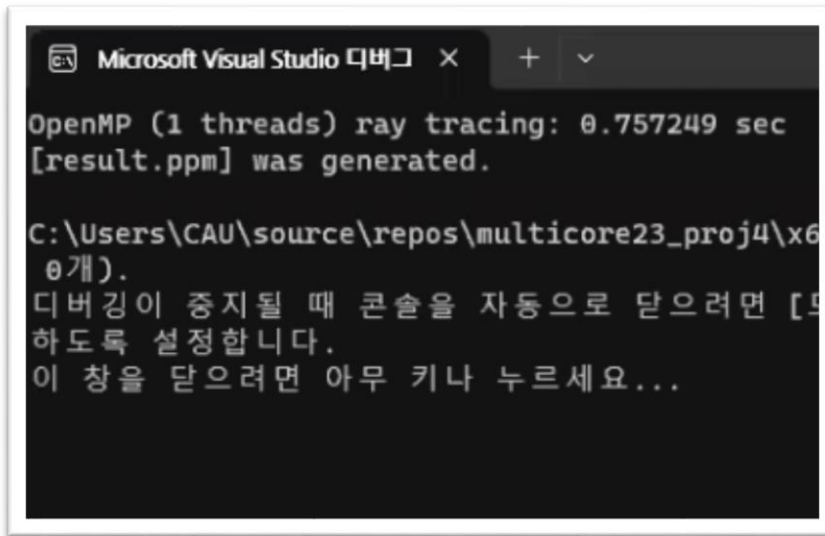
cout << "CUDA ray tracing: " << duration.count() / 1000000.0 << " sec" << endl;
cout << "[result.ppm] was generated." << endl;
return 0;
}

```

## [Results]

### - openmp\_ray.cpp


#### ■ 1 thread



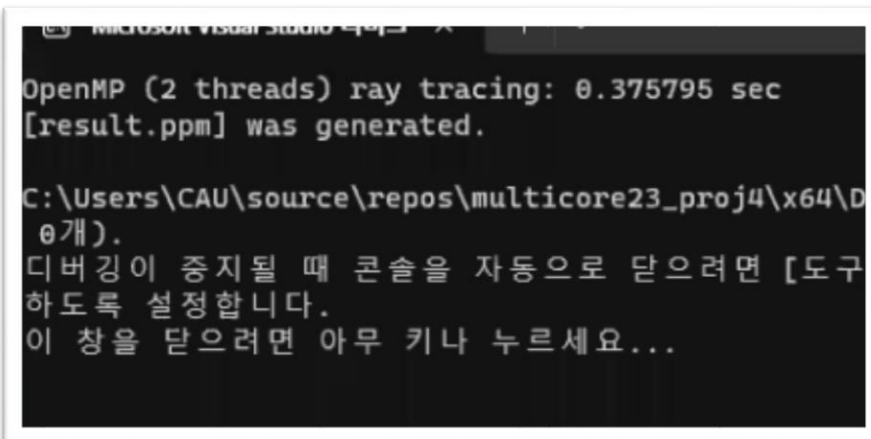
```
Microsoft Visual Studio 디버그 x + v

OpenMP (1 threads) ray tracing: 0.757249 sec
[result.ppm] was generated.

C:\Users\CAU\source\repos\multicore23_proj4\x64\Debug>
0개).
디버깅이 중지될 때 콘솔을 자동으로 닫으려면 [도구] 메뉴의
하도록 설정합니다.
이 창을 닫으려면 아무 키나 누르세요...
```




#### ■ 2 threads



```
Microsoft Visual Studio 디버그 x + v

OpenMP (2 threads) ray tracing: 0.375795 sec
[result.ppm] was generated.


C:\Users\CAU\source\repos\multicore23_proj4\x64\Debug>
0개).
디버깅이 중지될 때 콘솔을 자동으로 닫으려면 [도구] 메뉴의
하도록 설정합니다.
이 창을 닫으려면 아무 키나 누르세요...
```



■ 4 threads

```
Microsoft Visual Studio 디버그 x + v
OpenMP (4 threads) ray tracing: 0.212556 sec
[result.ppm] was generated.


C:\Users\CAU\source\repos\multicore23_proj4\...
0개).
디버깅이 중지될 때 콘솔을 자동으로 닫으려면 [
하도록 설정합니다.
이 창을 닫으려면 아무 키나 누르세요...
```



■ 6 threads

```
Microsoft Visual Studio 디버그 x + v
OpenMP (6 threads) ray tracing: 0.173832 sec
[result.ppm] was generated.


C:\Users\CAU\source\repos\multicore23_proj4\x6
0개).
디버깅이 중지될 때 콘솔을 자동으로 닫으려면 [도
하도록 설정합니다.
이 창을 닫으려면 아무 키나 누르세요...
```



■ 8 threads

```
Microsoft Visual Studio 디버그 x + v
OpenMP (8 threads) ray tracing: 0.128327 sec
[result.ppm] was generated.

C:\Users\CAU\source\repos\multicore23_proj4\x6
0개).
디버깅이 중지될 때 콘솔을 자동으로 닫으려면 [도
하도록 설정합니다.
이 창을 닫으려면 아무 키나 누르세요...
```




- 10 threads


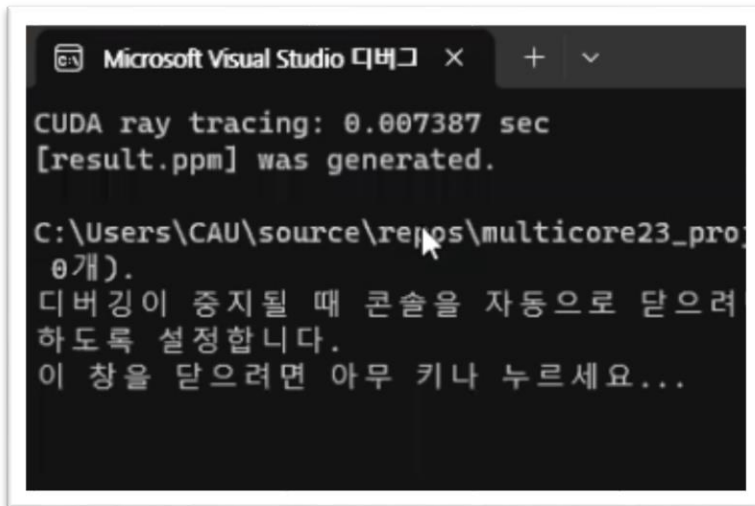
```

Microsoft Visual Studio 디버그 x + v
OpenMP (10 threads) ray tracing: 0.116590 sec
[result.ppm] was generated.

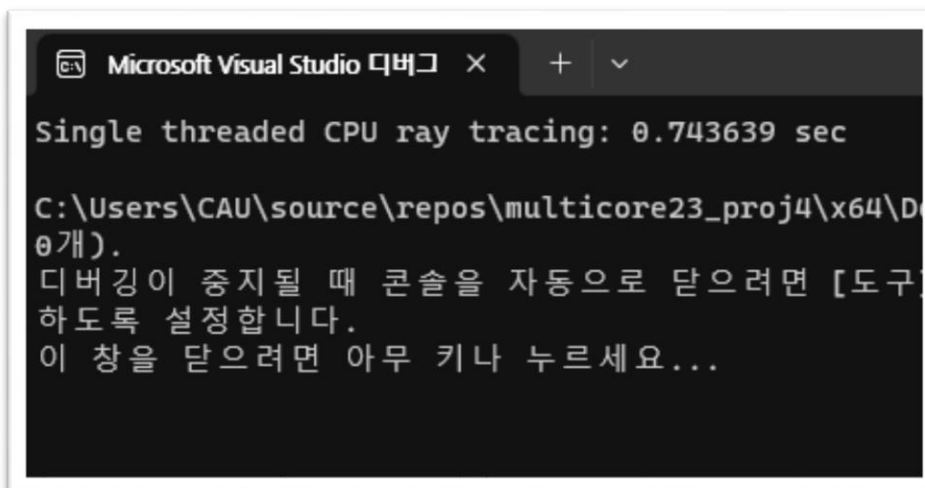
C:\Users\CAU\source\repos\multicore23_proj4\x64
0개).
디버깅이 중지될 때 콘솔을 자동으로 닫으려면 [도
하도록 설정합니다.
이 창을 닫으려면 아무 키나 누르세요...
  
```



- cuda\_ray.cu



- raytracing.cpp (Given source code)

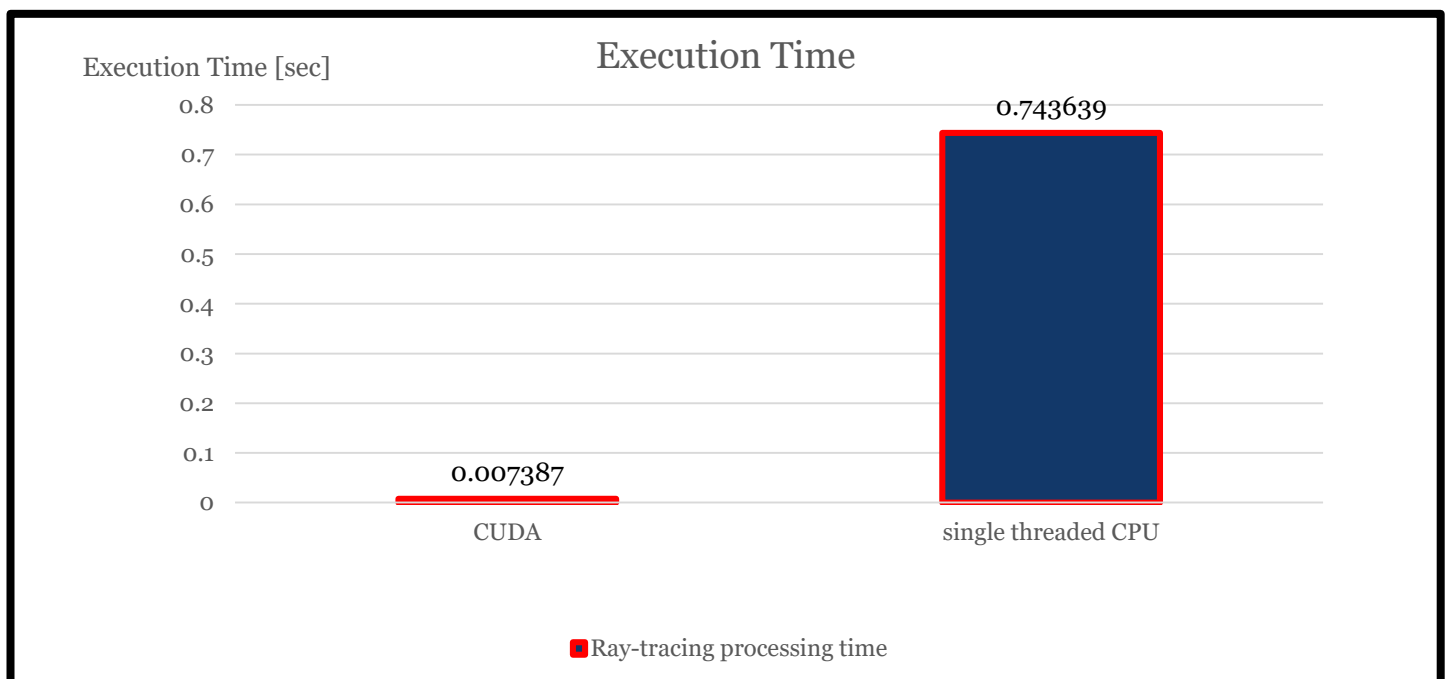



- Execution Time

OpenMp	1	2	4	6	8	10
Exec time	0.757249 sec	0.375795 sec	0.212556 sec	0.173832 sec	0.128327 sec	0.116590 sec

CUDA	
Exec time	0.007387 sec

Single threaded CPU	
Exec time	0.743639 sec







- **Performance**

OpenMp	1	2	4	6	8	10
Performance (1/exec time)	1.3205	2.6610	4.7046	5.7526	7.7925	8.5770

CUDA	
Performance (1/exec time)	153.3729

Single threaded CPU	
Performance (1/exec time)	1.3447



### **[Interpretation/Explanation on the Results]**

The above results shows that the more number of threads, the better performance. When the number of threads is small (ex. number of threads = 1), it takes 0.7 seconds. But, when the number of threads is large, program performance is improved dramatically. In other words, using OpenMp with large number of threads has good performance. Also, using CUDA library has much better performance than OpenMp. Using CUDA library is 20 times faster than OpenMp. It's because GPU has better performance in highly parallel applications than CPU.